



KubeCon



CloudNativeCon

S OPEN SOURCE SUMMIT

China 2019



Greg Kroah-Hartman



Disclaimer

- This talk vastly over-simplifies things.
- See notes for full details and resources.


<https://github.com/gregkh/presentation-mds>



- Same “family” of bugs as Spectre/Meltdown
- Hardware bugs
- Exploits the speculative execution model of modern CPUs.
- Many different variants.
- Is going to be with us for a very long time!




- MDS == “RIDL”, “Fallout”, “Zombieload”, and others
- CPU Hardware bugs
- Variants of the same basic problem
- Exploits the speculative execution model of Intel CPUs.
- Discovered by many different research teams
- Kernel and BIOS fixes required to fully solve




MDS

- One program can read another program's data
- Can cross the virtual machine boundary
- Exploits “hyper threading” (SMT) issues
- SMT are CPUs that usually share TLBs and L1 cache



OpenBSD was right

- Guessed more problems would be in this area
- Disabled SMT for Intel chips in June 2018
- Repeated the plea to disable this in August 2018
- Prevented almost all MDS issues automatically
- Security over performance
- Huge respect!



RIDL

- Rouge-Inflight-Data-Load
- Exploits CPU Line-fill buffers and Load ports
- Steal data across applications, virtual machines, secure enclaves
- Kernel fix by flushing CPU buffers/ports on context switch




Fallout

- Exploits CPU Store Buffers
- Read kernel data from userspace
- Breaks ASLR (random kernel addresses)
- “Meltdown” mitigation made this easier to exploit
- Kernel fix by flushing CPU buffers on context switch




Zombieland

- Exploits CPU Line-Fill buffers
- Much like RIDL
- Steal data across applications, virtual machines, secure enclaves
- Cool logo/name and demo
- Kernel fix by flushing CPU buffers on context switch




Other variants

- “Store-to-Leak forwarding”
- “Meltdown UC”
- All allow data to be stolen across security boundaries
- Kernel fix by flushing CPU buffers/ports on context switch



Flushing CPU buffers is slow

- All of these mitigations slow down the system
- No way yet to schedule different security domains on different physical processors (gang scheduling)
- Disabling SMT mitigates most problems (not ALL!)
- Must disable SMT and enable mitigations to solve completely.



Flushing CPU buffers is slow

- Performance numbers depend on your workload
- Kernel build
 - -2% smt=on
 - -15% smt=off
 - Heavily multi-threaded, CPU bound
- Kernel creation, no decrease
 - Single threaded, I/O bound
- syscalls are now expensive
- Test your workload!



Do you feel lucky?

- Users must now choose between performance and security
- What choice did your cloud provider choose?
- <https://make-linux-fast-again.com/>
 - Kernel builds faster by 15%!

Linux's response

- Kernel fixes available on announcement date
- Intel notified some kernel developers in advance
- Worked together across OS vendors to solve
- Much better than Spectre/Meltdown
- Process still needs to improve, Debian notified 48 hours before release.
- More fixes came after announcement
- Update your kernel and BIOS!

Linux security fixes

- Happen at least once a week
- Look like any other bugfix
- Rarely called out as security fix
- Many bugfixes not known to be security related until years later
- No differentiation between bug types
 - A bug is a bug is a bug
- Very few CVEs ever get assigned for kernel security issues

Linux security fixes != CVEs

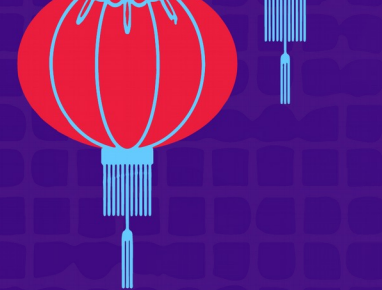

- Small fraction of kernel security fixes get CVEs
- If you only cherry-pick CVEs, you have an insecure system
- Some CVEs have follow-on fixes not documented anywhere
- [How the Linux Kernel Security team works](#)

Linux security fixes != CVEs

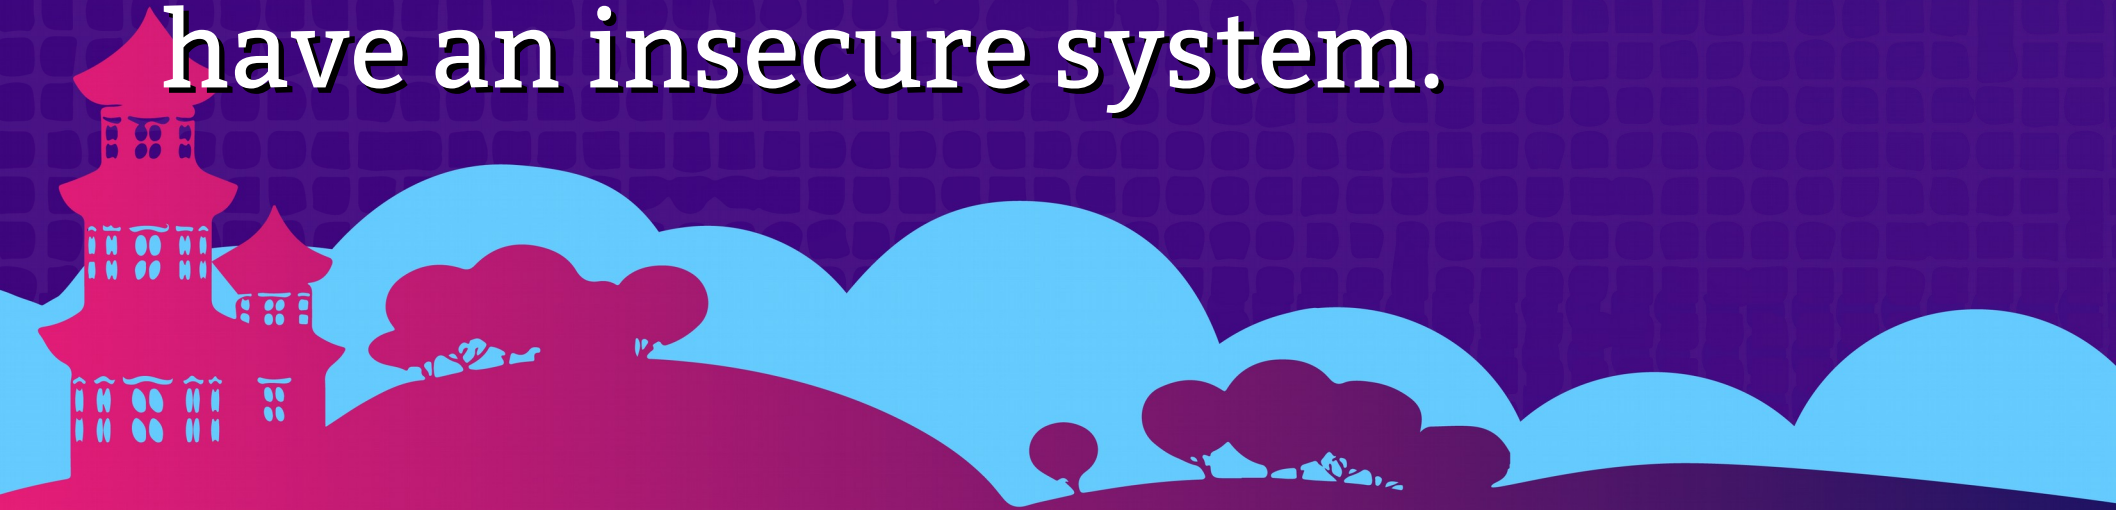
- Small fraction of kernel security fixes get CVEs
- 2006-2018 had 1005 CVEs assigned to the kernel
 - 41% (414) had a negative “fix date”
 - 12 never fixed
 - Average fix date, -100 days
 - Longest fix dates, -3897 and 2348 days
 - 88 fixed within 1 week
 - Standard deviation 405

Linux Longterm Kernels Fix Problems

- Bugs are fixed before you realize it is a issue.
- Google security team requests for Pixel phones in 2018:
 - 92% (201/218) problems were already fixed in LTS kernel
 - No need for cherry-picking or backporting
 - Remaining issues were due to out-of-tree code



If you are not using a supported Linux distribution kernel, or a stable / longterm kernel, you have an insecure system.



KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2019